

Common Statistical Models

VT-DSPG 2025

Anbin Rhee

July 2, 2025

Contents

1. What is a Model?
2. Simple Linear Regression
3. Multiple Linear Regression
4. ANOVA
5. Logistic Regression

What is a “Model”?

- Analogy: A toy airplane
 - A toy airplane is not a real plane, but it helps us understand how a real plane flies
- Statistical Models are similar
 - They are simplified mathematical representations of complex real-world data
 - They help us understand relationships between variables and make predictions



Key Concepts: Variables

- What is a variable?
 - Anything that can be measured or observed and can vary across observations
- Independent Variable (IV) : Predictor / Cause
 - The variable we change or observe to see if it has an effect
(The amount of fertilizer you give a plant)
- Dependent Variable (DV) : Outcome / Effect
 - The variable we measure to see if it changes in response to the independent variable
(The height of the plant)


Key Concepts: Variable Types

- Numerical
 - Can be any value within a range
 - Continuous : can take any values within a range
 - Discrete : can only take specific, separate values
 - e.g., temperature, sales, study hours, the number of children
- Categorical
 - Falls into distinct categories
 - e.g., gender, product type, teaching method (A/B/C)

Simple Linear Regression (SLR)

Example: Can we predict a student's Exam Score based on their Study Hours?

- We want to predict a continuous numerical outcome using one continuous numerical predictor



Simple Linear Regression

- The Conceptual Model (The 'True' Relationship in the population) :

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

- Y_i : The true outcome (Dependent variable) for individual i (e.g., Exam Score)
- X_i : The predictor (Independent variable) for individual i (e.g., Study Hours)
- β_0 : The true average value of Y when X is 0 (Intercept)
- β_1 : The true average change in Y for every one-unit increase in X (Slope)
- ϵ_i : The part of Y_i that the model cannot explain for individual i (Error term)

Simple Linear Regression

- The Estimated Model (From Our Data) :

$$\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_i$$

- \hat{Y}_i (Y-hat) : The predicted outcome for individual i (e.g., Exam Score)
- $\hat{\beta}_0$ (Estimated Intercept) : Our best guess for β_0 based on our data
- $\hat{\beta}_1$ (Estimated Slope) : Our best guess for β_1 based on our data
- Error vs. Residual
 - Error (ϵ_i) : Unobserved difference between the true Y_i and the true model's prediction
 - Residual (e_i) : Observed difference between the actual Y_i and our model's predicted \hat{Y}_i

$$e_i = Y_i - \hat{Y}_i : \text{We want small residuals!}$$

Simple Linear Regression: Assumptions

1. Independence of Errors

- Errors (ϵ_i) are independent of each other

2. Linearity

- The relationship between X and Y is approximately linear

3. Constant Variance of Errors (Homoscedasticity)

- The variability of the residuals should be roughly constant across all levels of the predictor X

4. Normality

- The residuals (e_i) are approximately normally distributed

Simple Linear Regression: Standard Error

- Standard Error (SE) of Coefficients:
 - How much our estimates would likely vary if we took many different samples from the same population
 - A smaller SE means our estimate is more precise and reliable (less-sample to sample variation)

Simple Linear Regression: T-Test

- Hypotheses

$$H_0: \beta_1 = 0 \text{ vs. } H_1: \beta_1 \neq 0$$

- We need statistical test used to get the p-value for individual coefficients in regression
- It calculates a t-statistics for each coefficient:

$$t = \frac{\text{Coefficient Estimate } (\beta_1) - \text{Hypothesized Value } (\beta_1 = 0 \text{ under } H_0)}{\text{Standard Error of Coefficient } (\beta_1)}$$

- A smaller SE means our estimate is more precise and reliable
- It measures how many standard errors away our estimated coefficient is from the hypothesized value
- A **large** |t| means our estimate is further from zero, making it **less** likely to be due to chance

SLR: Confidence Interval (CI)

- Instead of just a single point estimate, a CI provides a range of plausible values for the true population parameters
- Example: A 95% CI for the slope might be (a, b) . This means, “We are 95% confident that the true average change in Y for a one-unit change in X lies somewhere between a and b ”
- It gives us a sense of the precision of our estimate and the range of values we’d expect the true parameters to fall in if we repeated the experiment many times

Simple Linear Regression: Interpretation

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	37.6210	2.1316	17.65	<2e-16 ***
study_hours	5.5231	0.3314	16.66	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.13 on 48 degrees of freedom

Multiple R-squared: 0.8526, Adjusted R-squared: 0.8495

F-statistic: 277.7 on 1 and 48 DF, p-value: < 2.2e-16

- Intercept(37.6210): The predicted exam score for a student who studied 0 hours
- Study_hours (5.5231): For every additional hour of study, the exam score increases on average by 5.52 points

Multiple Linear Regression

Example: Can we predict an individual's Salary based on their Year of Experience, Education Level, and Hours worked Per Week?

- Dependent variable : Salary *Continuous*
- Independent variables : Year of Experience, *Continuous*

Education Level *Categorical*

Hours Worked per Week *Continuous*
- We want to predict/explain a **continuous numerical outcome** using two or more independent variables
- Extends simple linear regression by adding more predictors

Multiple Linear Regression

- The Conceptual Model (The 'True' Relationship in the population) :

$$Y_i = \beta_0 + \beta_1 X_{1i} + \beta_2 X_{2i} + \beta_3 X_{3i} + \epsilon_i$$

- It is adding more X terms to our simple linear regression equation
- Y_i : Salary for individual i
- X_{1i} : Years of Experience for individual i
- X_{2i} : Education Level for individual i (High school, Bachelor, Master, PhD)
- X_{3i} : Hours Worked per Week for individual i
- Each β_j (slope) now represents the average change in Y for a one-unit change in that specific X , while holding all other X variables constant.

Handling Categorical Predictors

- For handling categorical predictors, one category becomes the reference group
- Statistical software automatically convert categorical variables into ‘**dummy variables**’ (0s and 1s) to indicate each category
- The coefficients for other categories show the predicted difference compared to this reference group
- For the salary example, we can set High School as our reference group for Education Level

Multiple Linear Regression: Interpretation

```

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   35480.39   3987.61    8.898 4.05e-14 ***
experience     1907.66    85.39   22.341 < 2e-16 ***
education_levelBachelor  5375.48   1441.44    3.729 0.000328 ***
education_levelMaster  10936.71   1410.09    7.756 1.03e-11 ***
education_levelPhD    15381.78   1411.88   10.895 < 2e-16 ***
hours_per_week    115.86    93.38    1.241 0.217811
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4658 on 94 degrees of freedom
Multiple R-squared:  0.8648,    Adjusted R-squared:  0.8576
F-statistic: 120.3 on 5 and 94 DF,  p-value: < 2.2e-16

```

- Experience (1907.66): Each additional year of experience increases predicted salary by about \$1907, adjusting for education level and hours worked (very strong effect)
- Education level – Bachelor: \$5375 more than High School, on average

Comparing Multiple Groups: ANOVA

Example: The company wants to compare the average purchase amounts among customers exposed to three different marketing campaigns (A, B, and C) to see which campaign performs best.

- Comparing the means of a **continuous Dependent Variable (purchase amount)** among **three marketing groups (Campaign A, B, C)**
- Analysis Of Variance (ANOVA) helps us:
 - Determine if there is a statistically significant difference among the group means

Comparing Multiple Groups: ANOVA

- The Conceptual Model:

$$Y_i = \mu + \alpha_j + \epsilon_i$$

- Y_i : The outcome (Dependent variable) for individual i (Purchase Amount)
- μ : The overall average of Y across all groups (Grand Mean)
- α_j : The effect of being in group j (Campaign A, B, or C). It represents how much the mean of group j deviates from the grand mean
- ϵ_i : The part of Y_i that the model cannot explain

Comparing Multiple Groups: ANOVA

- Hypothesis:

$$H_0: \alpha_1 = \alpha_2 = \alpha_3 \text{ vs. } H_1: \text{at least one group differ}$$

- ANOVA looks at two types of variation in the data:
 1. Between-Group Variation: The variability among the group means themselves (effect)
 2. Within-Group Variation: The variability within each group (noise)

- F-statistics:

$$F = \frac{\text{Variability between Group Means}}{\text{Variability within Groups}}$$

- A larger 'F' value suggest the group differences are large compared to the random differences withing each group, indicating significant group effects

ANOVA: Interpretation

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	49.064	1.310	37.455	< 2e-16	***
campaignB	13.121	1.853	7.083	1.03e-11	***
campaignC	6.383	1.853	3.446	0.000652	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 13.1 on 297 degrees of freedom

Multiple R-squared: 0.1445, Adjusted R-squared: 0.1388

F-statistic: 25.09 on 2 and 297 DF, p-value: 8.556e-11

- CampaignB (13.121): Customers in Campaign B had an average purchase amount that was \$13.121 higher than customers in Campaign A. This difference is highly statistically significant

Logistic Regression

Example: The medical company has developed a new treatment to improve treatment success compared to the standard treatment, after accounting for patient age and gender

- Outcome (Y) is either success or failure *Binary*
- Predictors (X): Treatment group (New vs. Standard) *Categorical*
 - Age *Continuous*
 - Gender *Categorical*
- Linear regression predicts continuous values. It could predict “success” as -0.5 or 1.2, which don’t make sense for a Yes/No outcome

Logistic Regression

- Instead of predicting the outcome (0 or 1) directly, Logistic Regression predicts the **probability of the ‘Yes/Success’ outcome** occurring (e.g., the probability of treatment success)

- The Model:

$$\log(\text{Odds of success}) = \beta_0 + \beta_1 * \text{Treatment} + \beta_2 * \text{Age} + \beta_3 * \text{Gender}$$

- $$\text{Odds} = \frac{\text{Probability of Event}}{1 - \text{Probability of Event}} = \frac{\text{Probability of Success}}{\text{Probability of Failure}}$$
- Purpose: To predict the probability of a binary outcome and understand which factors increase or decrease the odds of that outcome

Logistic Regression

- The Model:

$$\log(\text{Odds of success}) = \beta_0 + \beta_1 * \text{Treatment} + \beta_2 * \text{Age} + \beta_3 * \text{Gender}$$

- $\text{Odds} = \frac{\text{Probability of Event}}{1 - \text{Probability of Event}}$
- β_1 : change in log-odds of success for new treatment compared to standard
- $\text{Exp}(\beta_1)$: **odds ratio of success** for new treatment
 - How many times higher or lower the odds of success are under the new treatment compared to the standard treatment
- Logistic regression models the **probability of success**, allowing adjustment for other factors (age, gender) while comparing treatment effects

Logistic Regression: Interpretation

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	1.01130	0.95185	1.062	0.2880
treatmentStandard	-0.74231	0.35994	-2.062	0.0392 *
age	-0.03402	0.01900	-1.790	0.0735 .
genderMale	-0.55005	0.35980	-1.529	0.1263

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

- TreatmentStandard (Estimate: -0.74231): Patients in the Standard treatment group have $\exp(-0.74231)=0.476$ times the odds (or are about 52.4% less likely) of achieving treatment success compared to patients in the New treatment group, holding age and gender constant. This difference is statistically significant.

Quick Quiz

- What statistical model would you use to compare the average productivity across these three training programs?

Example: A company tests three types of training programs to improve employee productivity. After the training, they measure each employee's number of tasks completed per day.

- **ANOVA model** to compare the mean productivity across three training program

Important Considerations

- Garbage in, garbage out
 - Bad data -> bad insight
- Correlation is NOT causation
 - Just because two variables are related does not mean one causes the other
- Models are Simplifications
 - They are mathematical abstractions of reality. They do not capture every nuance
- Overfitting
 - A model that is too complex for your data might fit your current data perfectly but perform poorly on new, unseen data

Summary

- **Simple Linear Regression:** Predicting a continuous outcome from one continuous predictor (e.g., Exam Score from Study Hours)
- **Multiple Linear Regression:** Predicting a continuous outcome from multiple predictors (continuous or categorical) (e.g., Salary from Experience, Education, Hours)
- **ANOVA:** Comparing the means of a continuous outcome across three or more groups (e.g., Purchase Amount from Campaign A, B, and C)
- **Logistic Regression:** Predicting a binary (Yes/No) outcome from one or more predictors (e.g., Treatment Success from Treatment Group, Age, Gender)